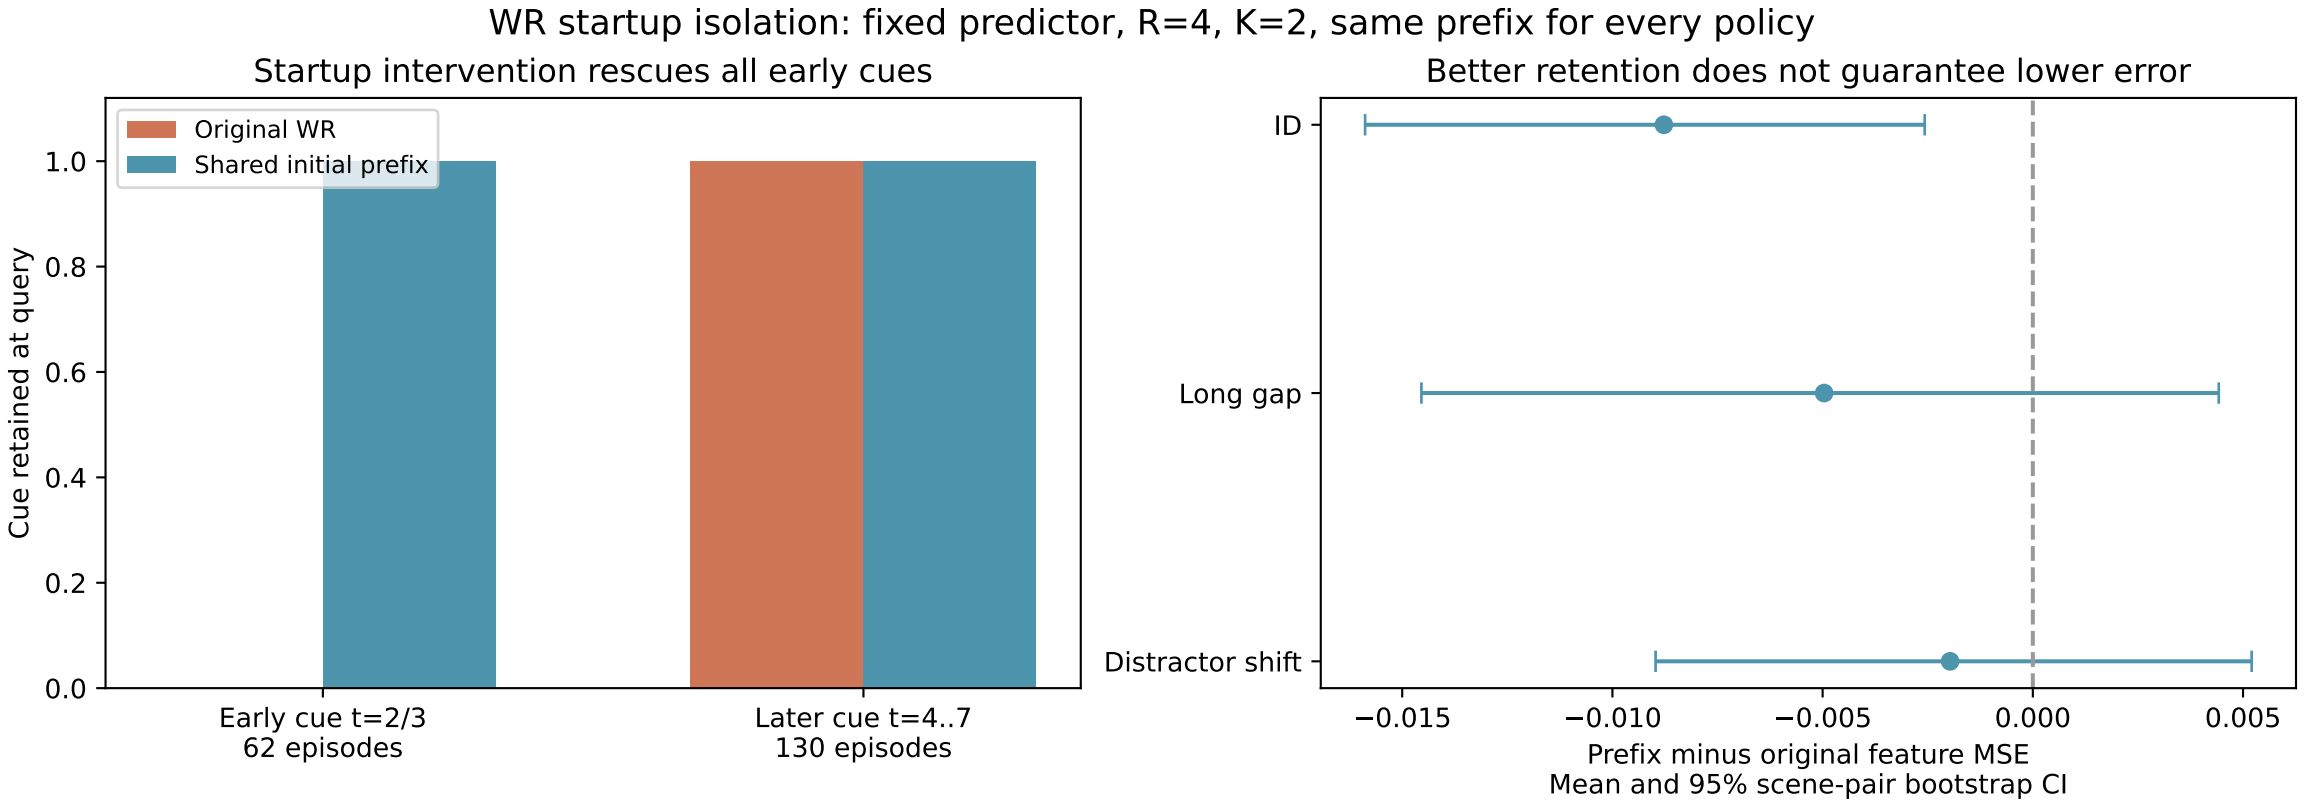
WR startup isolation: fixed predictor, R=4, K=2, same prefix for every policy
Startup intervention rescues all early cues
Original WR
Shared initial prefix
Cue retained at query
1.0
0.8
0.6
0.4
0.2
0.0
Early cue t=2/3
62 episodes
Later cue t=4..7
130 episodes
Better retention does not guarantee lower error
ID
Long gap
Distractor shift
−0.015
−0.010
−0.005
0.000
0.005
Prefix minus original feature MSE
Mean and 95% scene-pair bootstrap CI